



VIENNA DATA
SCIENCE TOOLS

NOVOMATIC

Our favorite Podcasts

- Software Engineering Daily
 - <https://softwareengineeringdaily.com/>
- This Week in Machine Learning & AI
 - <https://twimlai.com>
- O'Reilly Data Show
 - <https://www.oreilly.com/topics/oreilly-data-show-podcast>
- Software Engineering Radio (by IEEE)
 - <http://www.se-radio.net>
- The InfoQ Podcast
 - <https://www.infoq.com/the-infoq-podcast>



In a Nutshell

- FaunaDB is combining the scale and flexibility of NoSQL with the safety and data integrity of relational systems:
 - Transactional database system
 - ACID-compliant
 - Highly available
 - Distributed / horizontally scalable
 - Replication
 - Sharding
 - ACID-compliant distributed transactions
 - They call it "Relational NoSQL"
- Uses the Calvin architecture for distributed DBMS
 - [Thomson, Alexander et al. "Calvin: fast distributed transactions for partitioned database systems." SIGMOD Conference \(2012\).](#)

Calvin vs. Spanner

- Spanner
 - [Corbett, James C. et al. "Spanner: Google's Globally Distributed Database." ACM Trans. Comput. Syst. \(2012\).](#)
 - Spanner is Google's scalable, multi-version, globally- distributed, and synchronously replicated database.
 - It is the first system to distribute data at global scale and support externally-consistent distributed transactions.
 - Physical atomic clock HW synchronizes the system time on all shards within very small error bounds.
- Implementations of the Spanner architecture
 - CockroachDB
 - <https://www.cockroachlabs.com/>
 - YugaByte DB
 - <https://www.yugabyte.com/>
 - Apparently problematic due to lack of Google's HW infrastructure
 - <https://www.infoq.com/articles/relational-nosql-fauna>
 - <https://fauna.com/blog/distributed-consistency-at-scale-spanner-vs-calvin>
 - Counterargument from YugaByte
 - <https://blog.yugabyte.com/google-spanner-vs-calvin-global-consistency-at-scale/>

Thank you!

bpirvu@novomatic.com
jwilms@novomatic.com

NOVOMATIC